

Safety-helpfulness trade-off across model families

